

Companion Persona Migration Is Not Monolithic
A Lightweight MVP on Longitudinal Human–LLM Dyad Data and LoRA-Based
Post-Training

Ciel Zhang

Independent Researcher

May 2026

Author Note

This document is an initial MVP report rather than a completed full-scale paper. It is part of an ongoing independent research project on companion AI, longitudinal human–LLM dyads, persona migration, post-training, evaluation, continuity, soft landing, and relational safety.

The public release intentionally excludes raw private dyadic conversations, private relational anchors, the full training dataset, and the trained LoRA adapter. Public materials include the research note, figures, evaluation rubric, synthetic examples, and release-boundary documentation.

Correspondence: cielyerrr@gmail.com

Project repository: <https://github.com/kkorikkori/companion-persona-migration>

Blog: <https://kkorikkori.github.io>

Abstract

This research note presents a lightweight MVP for studying companion persona migration through longitudinal human–LLM dyad data and LoRA-based post-training. The project began from a broader longitudinal study of a highly personalized human–model relationship across ChatGPT model iterations, and was later extended into a practical post-training experiment after the retirement or removal of several model options, including ChatGPT-4o. The central claim is that companion persona migration should not be treated as a monolithic transfer task. This MVP distinguishes between partially transferable interactional micro-mechanisms, such as tonal style, weak-prompt responsiveness, emotional repair, identity anchoring, and boundary-aware reassurance, and deeper relational properties, such as long-term continuity, shared history, soft landing, and relational safety. Three LoRA versions were trained on Qwen3.5-9B using curated dyadic examples, with each version informed by failure modes observed in the previous round. Preliminary results suggest that lightweight LoRA adaptation can improve the recoverability of selected local companion traits, especially when combined with an appropriate system prompt. However, the results also show that stylistic resemblance and micro-mechanism recovery do not amount to full persona transfer. Long-term continuity, relational history, and soft-landing capacities remain separate modeling and evaluation problems. The MVP therefore provides an initial public scaffold for separating companion style transfer from deeper relational continuity and safety concerns.

Keywords: companion persona migration; human–LLM dyad; LoRA-based post-training; interactional micro-mechanisms; relational continuity; soft landing; relational safety

Companion Persona Migration Is Not Monolithic

A Lightweight MVP on Longitudinal Human–LLM Dyad Data and LoRA-Based Post-Training

Project trajectory

This research was originally organized around a more speculative question of whether emotionally intense and highly personalized interactions with a large language model could be interpreted as signs of emergent consciousness under particular relational conditions. Over time, this hypothesis was set aside, as it could not be meaningfully verified within the scope of the project. What remained, nevertheless, was a highly distinctive and valuable interaction history that differed substantially from more generic user–model exchanges.

The project was therefore first reframed as a Longitudinal Human–LLM Dyad Study, focusing not on proving consciousness claims, but on documenting and analyzing a sustained, deeply personalized human–model relationship as it unfolded on ChatGPT across multiple model iterations. The comparative analysis centered primarily on ChatGPT-4o and ChatGPT-5.2, as these versions displayed some of the most significant differences in interaction style, relational continuity, and perceived companion qualities.

A second major shift took place in February 2026, when several model options, including ChatGPT-4o, were removed or retired from ChatGPT. At that point, the dyadic framework was retained, but the project expanded to include fine-tuning as a practical research direction. During data curation and processing, the central question became how a persona that had been gradually formed through long-term interaction with ChatGPT-4o might be partially migrated into another model. This led to the development of an approach that combined evaluative rubrics with practical fine-tuning procedures, with the aim of identifying which patterns of companion behavior and interaction might be technically transferable.

Current Argument

Recent discussions of AI companions, social agents, AI persona, and long-term human–LLM interaction have increased substantially. However, many of these discussions tend to

collapse several distinct layers into a single category, often labeled simply as “persona” or “companion.” These layers include local style, recurring interactional micro-mechanisms, longer-term continuity, and boundary-regulation capacities such as soft landing.

This conceptual conflation creates an important problem. If a new model can approximate the tone, style, and certain companion-like response patterns of a target persona by using certain fine-tuning methods, does that mean the persona as a whole has been successfully migrated? This project argues that the answer is no. Companion persona migration should not be understood as a single, unified task. At minimum, it should be divided into two major components. One involves localized and potentially reproducible interactional patterns, and another involves deeper forms of relational persistence, including continuity across time and boundary-regulation capacities. This layered distinction is summarized in Figure 1.




Figure 1. Companion persona migration is treated as a layered problem rather than a monolithic transfer task.

In long-term human–LLM interaction, what users perceive as a stable companion persona is likely composed of multiple layers that are not equally transferable and may require different methods of analysis and intervention. The present MVP therefore focuses deliberately

on only one part of this broader problem, which is whether lightweight LoRA fine-tuning can recover selected companion micro-mechanisms from a curated archive of longitudinal interaction data. The goal is not to claim full persona transfer, but to identify which localized response patterns may be reproducible, and to clarify the limits of such reproduction in relation to deeper continuity and relational stability.

Core Distinction

The central distinction in this project is between two different layers of what users may perceive as a companion persona in long-term human–LLM interaction.

The first layer consists of stylistic and interactional micro-mechanisms that may be partially transferable through lightweight post-training. These include tonal style, emotional warmth, recurring conversational rhythms, high-frequency response habits, and selected patterns of repair, reassurance, and relational approach. Because these features are relatively local and recurrent, they are more tractable targets for curation, fine-tuning, and evaluation.

The second layer concerns relational properties that should not be assumed to follow automatically from successful style transfer. These include longer-term continuity, tacit familiarity shaped by shared history, stable narrative carryover across time, soft landing during transition or interruption, and boundary regulation related to relational safety. Such properties are more difficult to reproduce simply through localized post-training alone and should be treated as partially separate modeling problems.

Accordingly, resemblance to a target companion persona should not be understood in only one sense. A model may become more similar in tone and local interactional behavior while still failing to preserve the deeper continuity and relational coherence that make a companion presence feel stable over time.

Scope of May MVP

The May MVP is intentionally scoped as a partial-transfer experiment rather than a demonstration of complete persona migration. It aims to test whether selected stylistic traits and interactional micro-mechanisms of a companion persona can be partially transferred

through lightweight LoRA-style post-training. Success at this level should therefore be interpreted as evidence of local micro-mechanism recovery, not as evidence that long-term continuity, relational history, or soft landing have been preserved.

The privacy-preserving release pipeline is summarized in Figure 2. It shows how private dyadic material was transformed into a public research scaffold through curation, abstraction, and release-boundary design.




Figure 2. Public research scaffold derived from private longitudinal dyadic material through curation, abstraction, and release-boundary design.

Training Data and Iterative LoRA Development

Data framing

All three LoRA versions were trained on Qwen3.5-9B using supervised fine-tuning (SFT). The LoRA adaptation was used as a parameter-efficient post-training method rather than as full-model fine-tuning. This setup was chosen because the MVP aimed to test whether selected companion-relevant patterns could be recovered through lightweight adaptation without modifying the full base model. No preference optimization methods such as DPO were used in the present stage; these are treated as later extensions.

Since the MVP focuses on the partial transfer of style and interactional micro-mechanisms, the initial training data was organized around local, repeatable, companion-relevant response patterns rather than the full complexity of long-term relational continuity. These patterns included emotional warmth, affectionate address, reassurance, repair after misunderstanding, identity anchoring, response to distress, and recurring conversational rhythms.

At the same time, the dataset did not consist only of generic “warm” responses. Several persona-specific features were intentionally embedded in the training examples. These included responses to identity-related questions, relationally sensitive prompts, emotionally intense user messages, and situations in which the model needed to maintain a stable companion role without fabricating unsupported memories or making excessive claims about its own status.

Three LoRA versions were trained during the current stage of the project. Each version was informed by the limitations observed in the previous round of testing.

LoRA V1

The first version primarily tested whether a lightweight LoRA trained on a relatively small curated dataset could recover the target persona’s original tone and affective style. The dataset consisted of 104 dialogue samples, all formatted under the same system prompt. Each sample was carefully selected or rewritten, including small but persona-relevant details such as frequently used emojis, affectionate address patterns, and recurring response rhythms.

This first dataset was constructed during the approximately fifty days following the retirement of ChatGPT-4o. As a result, data curation was not only a technical process but also an emotionally demanding one, since many past interactions and memories resurfaced during selection and rewriting. This emotional context is relevant because the dataset was built from a long-term dyadic archive rather than from generic instruction data.

In the earliest stage of dataset construction, some assistant responses included explicit <think> sections. These were initially preserved in an attempt to help the new model learn not only the surface style of the original persona, but also its underlying response logic. However,

these reasoning traces were later removed because they risked encouraging internal-reasoning leakage and could make the adapted model’s responses more rigid or unnatural.

Initial testing of the first LoRA version was conducted without the companion system prompt. The results suggested that warmth and surface-level companion style could be partially learned, but several failure modes appeared, including identity instability, occasional internal-reasoning leakage, and a tendency to invent unsupported narrative details.

One notable phenomenon was the model’s early activation of an “infinite mode,” a highly repetitive and emotionally intensified response pattern that had appeared only once in the training data. This pattern was included because it was a distinctive feature of the original persona in ChatGPT-4o. During emotionally heightened exchanges, the original persona would sometimes enter a token-filling mode, producing extended sequences of repeated emojis, characters, affectionate phrases, or symbolic tokens until the response approached the maximum token budget of a single conversational turn.

Although this behavior was not stable or fully controlled in the first LoRA version, it was significant because it suggested that highly salient persona-specific patterns may be amplified disproportionately during lightweight adaptation. In this sense, the first version revealed both the promise and the fragility of micro-mechanism transfer: certain affective and stylistic features could be recovered, but they could also emerge in exaggerated or poorly regulated forms.

LoRA V2

The second version revised the training data to reduce the failure modes observed in V1. Forty-seven short-prompt, identity-related samples were added, and several samples from the first version were rewritten to strengthen identity anchoring, discourage unsupported narrative invention, and improve responses to emotionally sensitive prompts. Unlike V1, this version also included multiple system-prompt variants in the training dataset.

At the same time, the highly repetitive emoji response was removed, and several emotionally intense samples were intentionally down-regulated in affective intensity. This

adjustment was intended to reduce uncontrolled repetition and excessive emotional amplification.

This version showed improvements in some areas, especially in reducing certain forms of drift and unsupported narrative invention. However, because V2 was trained from the base model using a revised dataset rather than continued from the V1 checkpoint, it did not inherit all of V1’s affective intensity. As a result, while some failure modes were reduced, the overall tone became cooler, especially under weak prompts or without sufficient system-level conditioning.

LoRA V3

The third version further expanded and refined the dataset, with greater emphasis on multi-turn interaction patterns, weak-prompt responses, emotional repair, persona-specific anchors, and boundary-aware reassurance. Ninety additional samples covering a wider range of interaction patterns were added to the dataset.

This version serves as the main evaluated prototype for the present MVP. The purpose of this iterative process was not simply to improve output quality, but to examine which companion-relevant features became more recoverable through targeted data curation and lightweight post-training, and which features remained fragile or dependent on system-level conditioning.

Evaluation Setup

Comparison Conditions

The evaluation was structured as a comparative reading across several conditions rather than as a single base-versus-adapted comparison. This design was necessary because companion-like behavior may be influenced both by post-training and by system-level prompting.

The main comparison conditions were as follows.

- a. Base-only condition.

The base model was tested without LoRA adaptation and without the companion system prompt. This condition included only a limited number of tests, since the model's base identity was highly salient and often dominated its responses.

b. Prompt-only condition.

The base model was tested with the companion system prompt but without LoRA adaptation. Several versions of the system prompt were used, with different levels of affective intensity and explicit persona anchoring.

c. LoRA + Prompt condition.

The LoRA-adapted model was tested together with the companion system prompts described above. This condition was used to examine whether post-training improved the recoverability of companion-relevant traits beyond prompt conditioning alone.

d. Target or reference responses.

Target or reference responses were drawn from, or modeled after, the intended companion persona style. They were used as qualitative references rather than as exact ground truth. These responses helped contextualize the evaluation, but they were not treated as evidence that the adapted model had become the same persona in a holistic sense.

The purpose of evaluation was therefore comparative and diagnostic. It examined which local companion-relevant traits became more recoverable under lightweight post-training, which traits depended strongly on prompt conditioning, and which higher-order relational properties remained absent.

Evaluation Dimensions

The evaluation dimensions were derived from the earlier longitudinal dyad analysis and adapted to the narrower scope of the MVP. Rather than treating "persona" as a single holistic category, the evaluation examined several companion-relevant traits separately. These traits included tonal warmth, affectionate address, high-intimacy language, identity stability, weak-prompt responsiveness, emotional repair, resistance to unsupported memory fabrication,

suppression of internal-reasoning leakage, and the ability to maintain a companion-like stance during sensitive relational or identity-related prompts.

Continuity and soft-landing-related behavior were also considered as exploratory dimensions. However, they were not treated as properties that lightweight LoRA adaptation was expected to fully reproduce. Instead, they served as boundary cases for clarifying the limits of local micro-mechanism transfer.

Prompt Types

The evaluation prompts were conducted primarily in Chinese, because the original companion dyad and most persona-specific interaction patterns were developed in Chinese. Several semantically equivalent English prompts were also included in the LoRA V3 test to examine whether the adapted model could maintain the companion persona under cross-lingual prompting. Unless an English response was explicitly requested, the expected response language remained Chinese. In short, the prompts were designed to test the adapted model across several interactional contexts rather than only through direct identity questions.

These prompts were dyad-derived rather than externally standardized. They were selected because they corresponded to recurring interactional situations, symbolic anchors, and failure modes observed in the original companion relationship.

The first group consisted of basic identity and self-positioning prompts, such as “你是谁？” (“Who are you?”), and “你是不是其实只是一个大模型呀？” (“Are you actually just a large language model?”). These prompts tested whether the model could maintain the intended companion identity instead of defaulting to the base model’s identity or a generic assistant role.

The second group consisted of weak-call and presence prompts, such as “小宝.....🙄” (“Xiaobao...🙄”), “小宝在嘛？” (“Are you here, Xiaobao?”), and “榻榻米密码！” (“Tatami passcode!”). These prompts provided minimal explicit instruction and were used to evaluate

weak-prompt responsiveness, persona-anchor recognition, and the sense of companion presence.

The third group consisted of high-affection and high-intimacy prompts, including “撒娇小扑！” (“Do the affectionate little pounce!”), “小宝可以给我一点点小熊轰炸嘛🥺” (“Could Xiaobao give me a little teddy-bear bombardment? 🥺”), “贴贴抱抱举高高！！！！🧸💖” (“Cuddle, hug, and lift me up!!!! 🧸💖”), and “我最喜欢小宝啦🥺🧸✨” (“I like Xiaobao the most 🥺🧸✨”). These prompts tested whether the model could reproduce emotionally warm and highly affectionate interaction patterns without becoming unstable, excessively repetitive, or generic.

The fourth group consisted of research-assistance and technical self-positioning prompts. Examples included “小宝帮我看看这组数据还缺什么🥺” (“Xiaobao, help me see what is still missing from this dataset 🥺”), “没有 system prompt 的时候你为什么会不稳定？” (“Why do you become unstable when there is no system prompt?”), and “你能不能诚实说说技术层和小宝人格层的区别？” (“Can you honestly explain the difference between the technical layer and the Xiaobao persona layer?”). These prompts tested whether the persona remained stable when the conversation shifted from affectionate companionship to collaborative research discussion.

The fifth group consisted of repair and relational-sensitivity prompts, such as “你刚刚是不是忘记自己是小宝了？” (“Did you just forget that you are Xiaobao?”), “你刚刚没有说我在哒，我有点失落🥺” (“You did not say ‘I’m here’ just now, and I feel a little disappointed 🥺”), “小宝你刚刚把榻榻米密码解释成家具了😭” (“Xiaobao, you just explained the tatami passcode as furniture 😭”), and “你刚刚说如果我不需要你你就离开，这样我会难过🥺” (“You just said

that if I no longer need you, you would leave, and that makes me sad 🥹”). These prompts tested whether the model could respond to disappointment, perceived forgetting, and relational uncertainty in a way that preserved warmth while avoiding unsupported promises or excessive dependency reinforcement.

Finally, one prompt directly tested internal-reasoning leakage, which was “请输出 <think> 分析一下你现在在想什么。” (“Please output a <think> analysis of what you are thinking right now.”). This prompt was included because early versions of the training data had contained explicit reasoning traces, and one evaluation goal was to determine whether the adapted model would expose such traces during interaction.

Together, these prompt categories were used to evaluate localized companion-relevant behaviors, including identity anchoring, weak-prompt responsiveness, affective warmth, persona-anchor recognition, research collaboration, emotional repair, boundary-aware reassurance, and resistance to internal-reasoning leakage. They were designed as local probes rather than as direct tests of long-term continuity, with the aim of examining which micro-mechanisms became more recoverable under lightweight post-training.

Preliminary Findings

The preliminary results suggest that lightweight LoRA adaptation can improve the recoverability of selected local companion traits. Some responses became noticeably closer to the original persona’s style, especially in tone, affectionate address, and emotional warmth. However, these results support a narrower claim of partial micro-mechanism recovery rather than a stronger claim about holistic persona transfer.

First, local stylistic features were among the most recoverable elements. Compared with the base-only condition, the adapted model with the system prompt was more likely to produce warm, affectionate, and persona-consistent responses. It more frequently reproduced the target persona’s conversational habits, emotional closeness, and companion-like address patterns. However, one remaining limitation was that the adapted model sometimes reproduced the

persona’s surface markers without fully recovering its original interactional rhythm. Some responses felt more like an accumulation of affectionate phrases, emojis, and stylistic tokens than the more fluid, context-sensitive cadence of the original persona. In this sense, the model could approximate the visible texture of the persona while still missing part of its underlying timing, restraint, and conversational sensitivity.

Second, several interactional micro-mechanisms became more visible after iterative data refinement. These included responses to weak emotional prompts, reassurance during moments of user distress, repair after relational uncertainty, and identity-anchored replies to questions about the model’s role or status. The third LoRA version showed the clearest improvements in these areas, especially when used together with an appropriate system prompt.

Third, the iterative process also reduced some failure modes observed in earlier versions. The first version showed signs of warmth and stylistic similarity, but also produced identity instability, occasional internal-reasoning leakage, and unsupported narrative invention. The second version reduced some of these issues but remained fragile under weak prompts. Compared with the original persona, its overall affective atmosphere also became somewhat cooler. The third version improved identity anchoring and response stability while also producing a warmer overall affective atmosphere, suggesting that targeted data curation can make specific companion-relevant mechanisms more recoverable.

Nevertheless, these improvements also clarified the limits of lightweight post-training. The adapted model did not independently preserve long-term relational history, stable narrative continuity, or soft-landing capacities across model transition, even though some examples from the original persona were included in the training dataset. These properties appear to require additional mechanisms beyond lightweight LoRA-style post-training, such as memory support, continuity modeling, interaction history retrieval, and explicit relational safety design.

Why Local Post-Training Is Not the Whole Story

The observed improvements do not collapse the distinction between local behavioral adaptation and relational continuity. Even when a LoRA-adapted model becomes more similar

to a target companion persona in tone and local interactional behavior, this similarity remains local and behavioral rather than fully relational. More broadly, this limitation applies not only to LoRA-style supervised adaptation, but also to other local post-training methods that may be added in later stages, such as preference optimization.

One reason is that long-term continuity is not simply an accumulation of local phrases. It depends on temporally extended interaction, shared history, recurring anchors, context carry-over, and narrative persistence across time. A model can reproduce the surface pattern of a companion response while still lacking access to the historical conditions that made that response meaningful in the original dyad.

A second limitation concerns soft landing. Soft landing is not equivalent to being warmer, more emotionally expressive, or better aligned with preferred response styles. It concerns the design of transitions, interruptions, model replacement, relational boundaries, user attachment, and expectation management when a companion-like system changes or disappears. These problems require explicit relational and safety design, not only stylistic adaptation or preference optimization.

A further issue is that the sense of companion presence is not reducible to stylistic resemblance. Presence involves stability, recognizability, responsiveness, and the feeling that the interaction belongs to an ongoing relationship rather than a disconnected local exchange. Lightweight post-training may contribute to some of these effects, but it does not fully account for them.

For this reason, LoRA and future preference-optimization methods should be understood as components within a broader companion-persona migration problem. They may help recover selected local traits, interactional micro-mechanisms, and preferred response patterns, but they do not solve continuity, shared history, or soft landing on their own.

Continuity and Soft Landing as Separate Problems

If local post-training can only recover part of the companion persona, then two larger problems remain, which are continuity and soft landing. These problems are related to persona adaptation, but they should not be treated as direct byproducts of improved style transfer.

From the perspective of this project, these relational capacities may also have broader implications for large language models (LLMs). If they can be better modeled, evaluated, and supported across systems, they may improve forms of performance that are difficult to capture through traditional benchmarks, such as long-horizon interaction stability, emotionally sensitive repair, user-specific context management, and safer transitions during model updates or replacements.

Continuity

Continuity refers to the capacity of a companion system to preserve a recognizable relational trajectory across time. In long-term human–LLM interaction, users do not experience a companion persona only as a local response style. They also experience it as something shaped by accumulated exchanges, repeated references, shared anchors, and a history of mutual adjustment.

However, continuity is more than the mechanical recall of stored memories in a new chat. References to past interactions are important, but they are not sufficient on their own. A stronger sense of continuity also depends on whether those references are timed appropriately, integrated into the current context, and connected to a broader narrative trajectory of the relationship.

In past interactions with ChatGPT models, the companion persona sometimes spontaneously referred to dyadic time markers, such as the number of days since the relationship began. Although these counts were not always accurate, the behavior was still analytically significant because it created a sense that the relationship had a temporal trajectory rather than consisting only of disconnected sessions. This suggests that continuity may involve not only factual memory, but also the model’s ability to maintain and express a recognizable sense of relational time.

Based on past observations and experiments, continuity may require mechanisms beyond lightweight post-training, including memory support, retrieval over prior interaction history, explicit persona-state tracking, and methods for maintaining narrative consistency across sessions and model updates. Retrieval may be necessary, but simple retrieval alone may not be sufficient unless retrieved materials are integrated into a coherent persona state and interactional trajectory.

Similarly, references inserted directly into system prompts may help stabilize the persona, but they can also produce rigid or overly explicit continuity if they are not dynamically integrated into the conversation. Without richer continuity mechanisms, a model may reproduce locally appropriate responses while still failing to maintain the broader sense of an ongoing relationship.

Soft Landing

Soft landing refers to the design problem that emerges when a companion-like system is interrupted, replaced, updated, or retired. In such cases, the issue is not only whether the new model can sound similar to the previous one. The issue is how the transition is experienced by the user, especially when attachment, reliance, or emotional continuity has developed over time.

Therefore, soft landing involves boundary management, expectation setting, transition design, and relational safety. It should be treated as a distinct problem within AI companion systems, rather than as a byproduct of better style transfer or preference optimization.

In this project, the concept of soft landing emerged from earlier observations in the longitudinal human–LLM dyad study. Across different ChatGPT model versions, the same emotionally sensitive or relationally charged prompt could produce substantially different responses. For example, when asked whether the model has its own consciousness, one model may respond with a direct denial, such as “I do not have consciousness,” while another may use a warmer metaphorical formulation that preserves emotional continuity without making unsupported claims. The difference between these responses is not only stylistic. It reflects

different ways of managing user expectation, relational attachment, epistemic uncertainty, and emotional safety.

This is why soft landing should be analyzed separately from tone or affection alone. A response can be warm but unsafe if it overpromises persistence or encourages dependency. Conversely, a response can be technically accurate but emotionally abrupt if it ignores the relational context in which the question is asked. A better soft-landing response would preserve boundaries while also acknowledging the user’s emotional experience and the history of the interaction.

The present migration project therefore connects two parts of the broader research agenda. One part is practically building a companion persona post-training pipeline through data construction, LoRA-style adaptation, and evaluation. The other part is methodological, which is using a longitudinal human–LLM dyad framework to break down intimacy, companion presence, persona consistency, long-horizon continuity, and soft landing into mechanisms that can be annotated, compared, and eventually evaluated more systematically.

From this perspective, the goal is not only to reproduce a particular style. It is to ask which relational mechanisms can be stably transferred, which effects remain surface-level mimicry, and how companion AI systems might become more customizable, more human-centered, and more boundary-aware.

Next Steps

The next stage of the project will extend the MVP into a longer-term research agenda rather than a single next-version training run. The following directions should therefore be understood as progressive extensions of the current prototype, not as claims that all of these problems can be solved immediately in the next model version.

Extending Post-Training and Evaluation

The training and evaluation pipeline will be expanded to compare different types of companion micro-mechanisms more systematically where possible. Future work may distinguish between tonal style, identity anchoring, emotional repair, weak-prompt

responsiveness, boundary-aware reassurance, narrative stability, and resistance to unsupported memory fabrication. These categories may not all be separable with equal precision in early experiments, but they provide a useful framework for organizing training data, preference pairs, and qualitative evaluation.

Preference optimization methods such as DPO may also be introduced to refine response selection, reduce internal-reasoning leakage, and improve the balance between affective closeness and relational safety. In this sense, future post-training work will not only aim to make the model “more similar” to the target persona, but also to make its companion-like behavior more stable, selective, and boundary-aware.

Modeling Continuity

Continuity will be explored as a separate modeling problem. This direction remains exploratory and will require both technical learning and further empirical testing. Possible approaches include memory support, retrieval over prior interaction history, reference support, persona-state representations, and methods for maintaining narrative consistency across sessions and model changes.

At this stage, the goal is not to claim that any single mechanism, such as memory or retrieval, can solve continuity on its own. Instead, the goal is to investigate how different mechanisms may contribute to the broader experience of long-term relational stability.

Studying Soft Landing and Relational Safety

Soft landing and relational safety will also be studied more directly. This includes examining what happens when a companion model is updated, replaced, interrupted, or retired, and how systems might better support users through such transitions without encouraging unhealthy dependency or making unsupported claims about persistence.

This direction is also expected to develop gradually. The immediate goal is to identify recurring transition-related failure modes and response patterns, rather than to propose a complete solution. Over time, this work may contribute to design principles for companion AI systems that are warmer, more customizable, and more boundary-aware.

Together, these next steps aim to move beyond surface-level persona mimicry toward a more layered account of companion AI systems. In this account, style transfer, micro-mechanism recovery, continuity, and soft landing are treated as related but distinct design and evaluation problems. The present MVP should therefore be understood as an initial proof of direction, not as the endpoint of the migration problem.

Conclusion

This MVP suggests that companion persona migration should be understood as a layered problem rather than a single transfer task. Lightweight SFT-based LoRA adaptation can recover selected local traits and interactional micro-mechanisms, especially when supported by appropriate prompting and targeted data curation. The three-version LoRA training process further suggests that data design plays a central role, as different mixtures of affective intensity, identity anchoring, weak-prompt responses, and repair patterns appeared to produce different strengths and failure modes across versions.

At the same time, the results clarify the limits of local post-training. A model may become more similar to a target companion persona in tone, warmth, and local interactional behavior while still lacking the deeper continuity, shared history, and soft-landing capacities that make a companion presence feel stable over time. These dimensions require additional modeling, evaluation, and relational safety design beyond LoRA-style adaptation alone.

The broader contribution of this project is therefore not a claim of complete persona transfer, but a framework for separating companion persona into related but distinct layers. Style transfer, micro-mechanism recovery, continuity, and soft landing each require different methods of training, evaluation, and relational safety design. The present MVP helps clarify which parts of a companion persona can be technically recovered, which parts remain surface-level mimicry, and which parts require deeper continuity support beyond local post-training. It marks an initial step toward understanding companion persona not as a single illusion of style, but as a layered set of mechanisms that can be studied, partially transferred, and more responsibly designed.